



Computer Vision (Course Code: 4047)

Module-4:Lecture-3: Deep Learning Classifiers - basics

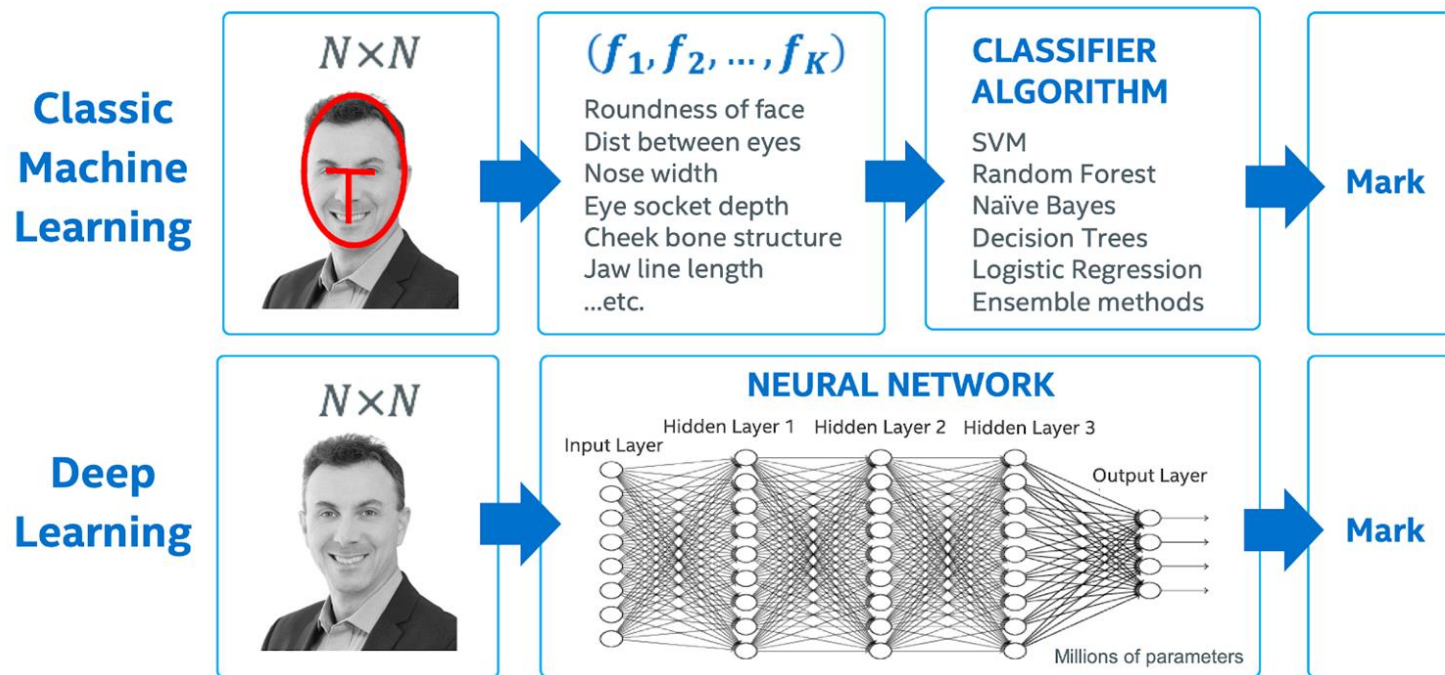
Gundimeda Venugopal, Professor of Practice, SCOPE

Topic Coverage

- Deep Learning
- Why Deep Learning?
- Deep Neural Network Optimization
- Convolutional Neural Network
- CNN Usecases

Deep Learning

- Deep Learning is a subset of Machine Learning based on Neural Networks that permit a machine to train itself to perform a task.
- Deep learning algorithms attempt to learn (multiple levels of) representation by using a hierarchy of multiple layers (> 1 hidden layer).
- Each layer of nodes trains on a distinct set of features based on the previous layer's output
- The further you advance into the neural net, the more complex the features your nodes can recognize, since they aggregate and recombine features from the previous layer



- Manually designed features
- Need small amount of training data
- Very good at learning patterns. Small models
- Cannot model complex challenges with high accuracy

- Model complex challenges (Vision, speech, NLP, ..)
- Automatic Feature Learning. Easy to adapt, fast to learn
- Utilize large amounts of training data
- Requires better hardware resources for training

Face Detection example

Deep Learning Pioneers

2018 ACM A.M. TURING AWARD RECIPIENTS



Geoffrey Hinton



Yoshua Bengio



Yann LeCun

- Geoffrey Hinton: Backpropagation, Boltzman machines, Deep Learning, CNN improvements, Capsules, AlexNet (Google)
- Yann LeCun: CNN, Backpropagation improvements (Facebook)
- Joshua Bengio: GAN, word embeddings, sequence models, (Mila)

[Andrew Ng](#) (Google Brain, Baidu)

[Ian Goodfellow](#) (Apple, GAN)

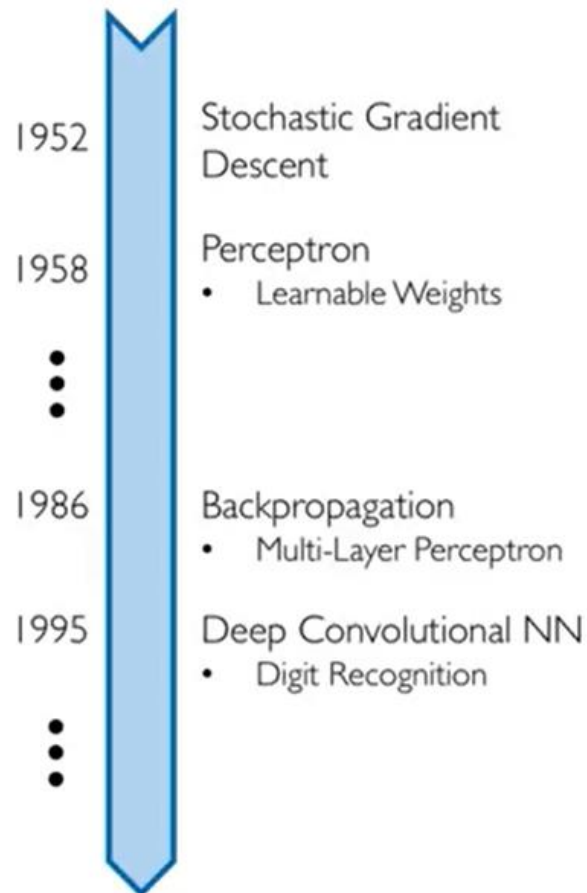
[Fei-Fei Li](#) (Stanford Prof, Google)

[Andrej Karpathy](#) (Tesla, ImageNet)

[Demis Hassabis](#) (Deep Mind)

Many others ..

Why Now?



Neural Networks date back decades, so why the resurgence?

1. Big Data

- Larger Datasets
- Easier Collection & Storage

IMAGENET



WIKIPEDIA
The Free Encyclopedia



2. Hardware

- Graphics Processing Units (GPUs)
- Massively Parallelizable



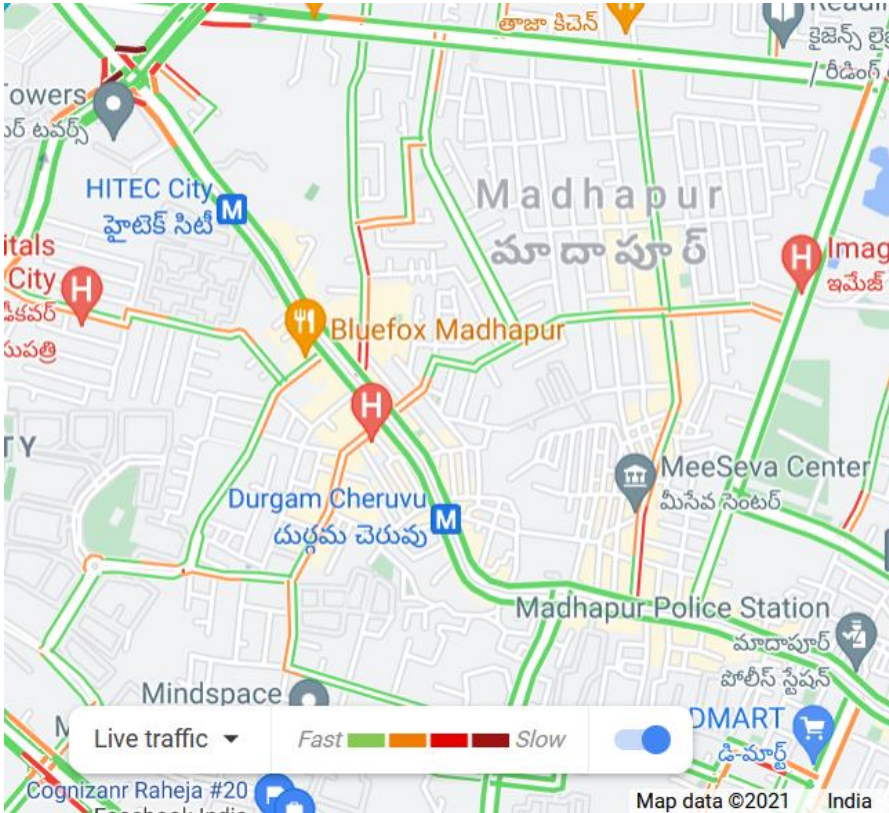
3. Software

- Improved Techniques
- New Models
- Toolboxes



Real world examples

Google Maps Traffic Prediction




Self driving Cars



PERCEPTION



LOCALIZATION



PLANNING



CONTROL

In **Perception**, you find the **environment** and **obstacles** around you.

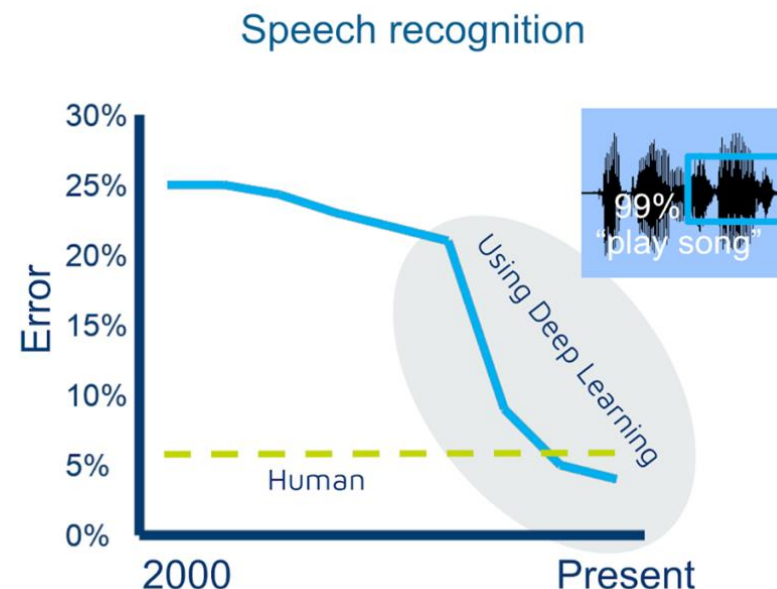

In **Localization**, you define your **position** in the world at 1–3 cm accuracy.

In **Planning**, you define a **trajectory** from **A** to **B**, using perception and localization.

In **Control**, you follow the trajectory by **generating a steering angle** and an **acceleration value**.

State of the art in Image Recognition, Speech and NLP

In ~2010 Deep Learning started outperforming other Machine Learning techniques first in speech and vision, then NLP




Several big improvements in recent years in NLP

- Machine Translation
- Sentiment Analysis
- Dialogue Agents
- Question Answering
- Text Classification ...

Leverage different levels of representation


- words & characters
- syntax & semantics

Neuron




Artificial neuron is inspired by biological neuron

Activation function



Train a neural network




- Choose network design
- Form a neural network
- Choose hyper parameters
- Compute an estimate value for all samples
- Compute loss
- Reduce loss
- Repeat last three steps

Error and Loss function

- In most learning networks, error is calculated as the difference between the actual output and the predicted output.
- The function that is used to compute this error is known as Loss Function.
- Different loss functions will give different errors for the same prediction, and thus have a considerable effect on the performance of the model.
- One of the most widely used loss function is mean square error, which calculates the square of difference between actual value and predicted value.
- Different loss functions are used to deal with different type of tasks, i.e. regression and classification.


Neural Network Training




1. Need to optimize (min. Or max.) the objective/cost function $J(W)$ using an optimizer (e.g, Gradient Descent, SGD, RMSprop, Adam)
2. Generate error signal that measures difference between predictions and target values
3. Use error signal to change the weights and get more accurate predictions
4. Subtracting a fraction of the gradient moves you towards the (local) minimum of the cost function

Loss Functions in real life are quite complex. They may have:

- A Global Minima
- Multiple Local Minima
- Saddle points




Gradient



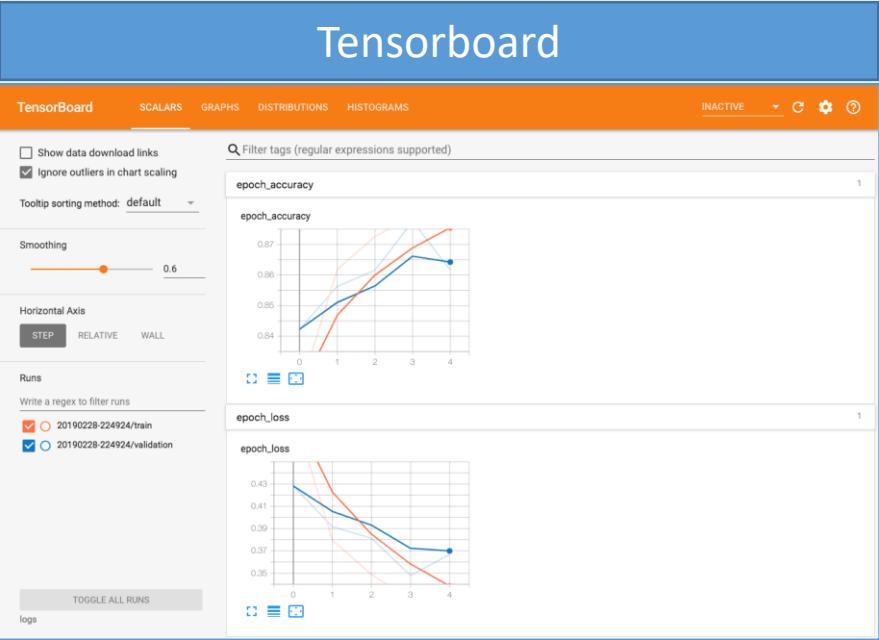
Optimisation functions usually calculate the **gradient** i.e. the partial derivative of loss function with respect to weights, and the weights are modified in the opposite direction of the calculated gradient. This cycle is repeated until we reach the minima of loss function.

Gradient Descent



The procedure of repeatedly evaluating the gradient and then performing a parameter update is called Gradient Descent.

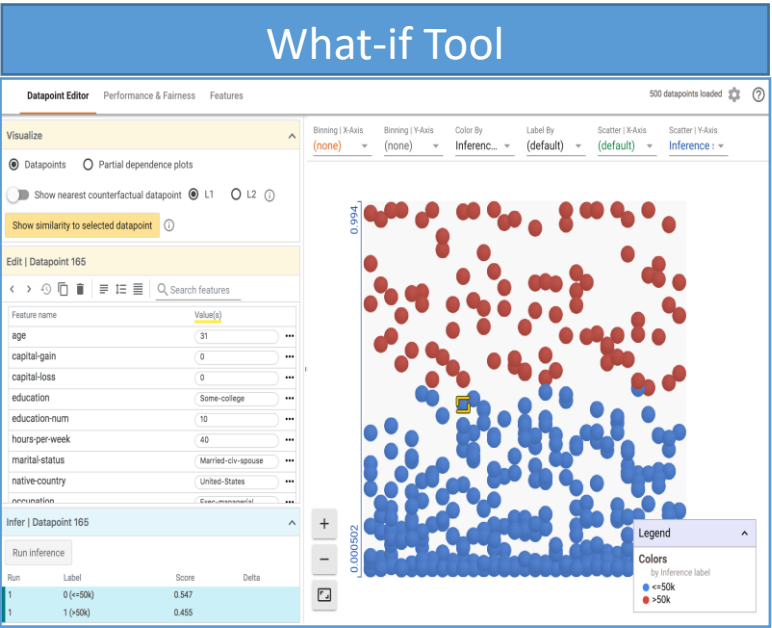
Google Tools for Better Model Building



Source: [Google](#)

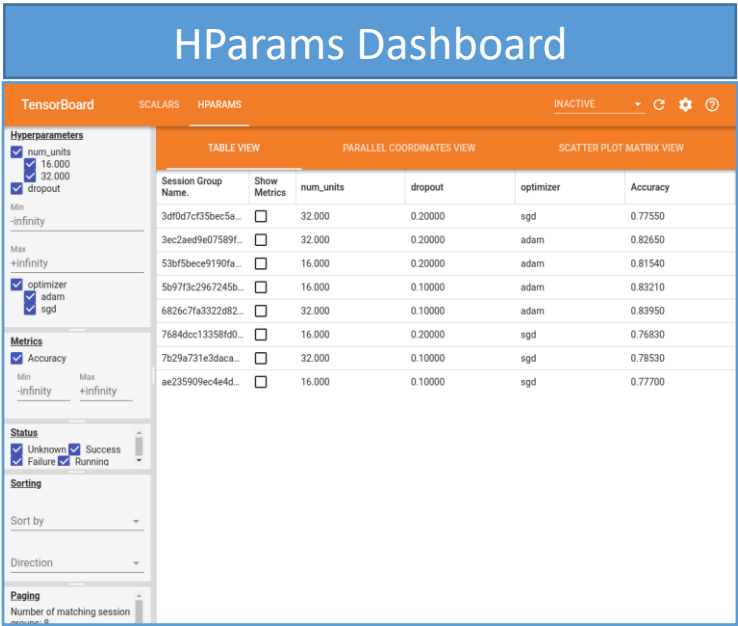
A Visual Interface to visualize the Model Graph and Training progress. Integrated with what-if and HPARAMs dash board tools.

Source: [Google](#)



Source: [Google](#)

A Visual Interface that helps to better understand your data sets and the output of machine learning models (Currently, Classification and regression models)



Source: [google](#)

HPARAMs dashboard is a tool for Hyperparameter tuning. It helps with this process of identifying the best experiment or most promising sets of hyperparameters.


Loss function Optimization - Learning Rate

Optimization through gradient descent


$$W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$$

How can we set the learning rate?


Stable learning rates converge smoothly and avoid local minima



Small learning rate converges slowly and gets stuck in false local minima



Large learning rates overshoot, become unstable and diverge



Learning rate

- With low learning rate the improvements will be linear
- Higher learning rate can decay the loss faster, but it can get stuck
- Initially keep the learning rate higher
- Later decrease the learning rate


Convolutional Neural Networks

Convolutional Neural Network


- Convolutional Neural Networks (CNNs) are Neural Networks that are designed to work efficiently for Images. Also, CNNs are used for NLP, Speech recognition and other tasks too.
- Few real world examples include:



Fingerprint Recognition



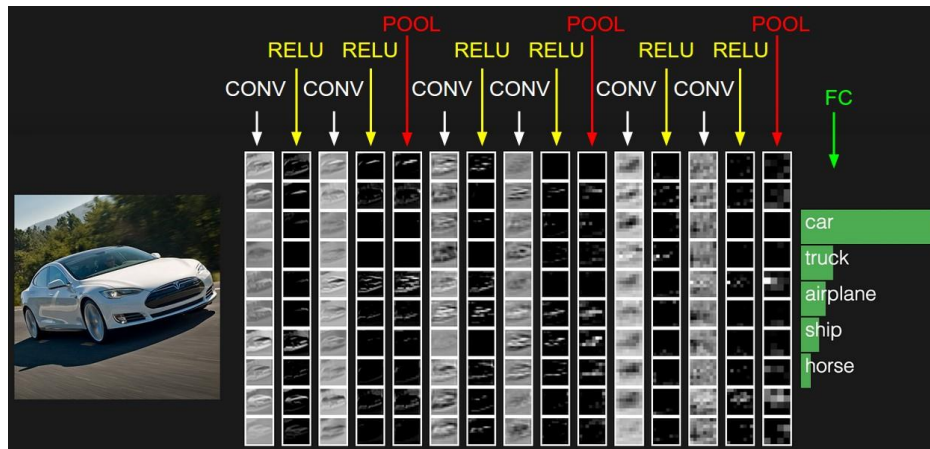
Face Recognition



Diabetic Retinopathy Detection

Convolutional Neural Network Structure




- Convolutional Neural Networks (CNNs) are Neural Networks that are designed to work efficiently for Images.
- Images can be Black/white, Grayscale and RGB (3 channels).
- A CNN is made up of Layers. Unlike a regular Neural Network, the layers of a CNN have neurons arranged in 3 dimensions: width, height, depth. The depth indicates the number of channels in the input image.
- A CNN architecture is a list of Layers that transform the image volume into an output volume (e.g. holding the class scores)
e.g., 32 X 32 X 3 input CIFAR-10 image 1 X 1 X 10 size output (10 output classes to choose)
- Multiple types of layers are used to build ConvNet architectures. We will stack these layers to form a full CNN architecture.
 - Convolutional Layer (CONV)
 - Pooling Layer (POOL)
 - Fully-Connected Layer (FC Layer)
 - Activation Layer (e.g., RELU layer)
- Each Layer accepts an input 3D volume and transforms it to an output 3D volume through a differentiable function
- Each Layer may or may not have parameters (e.g. CONV/FC do, RELU/POOL don't)
- Each Layer may or may not have additional hyperparameters (e.g. CONV/FC/POOL do, RELU doesn't)





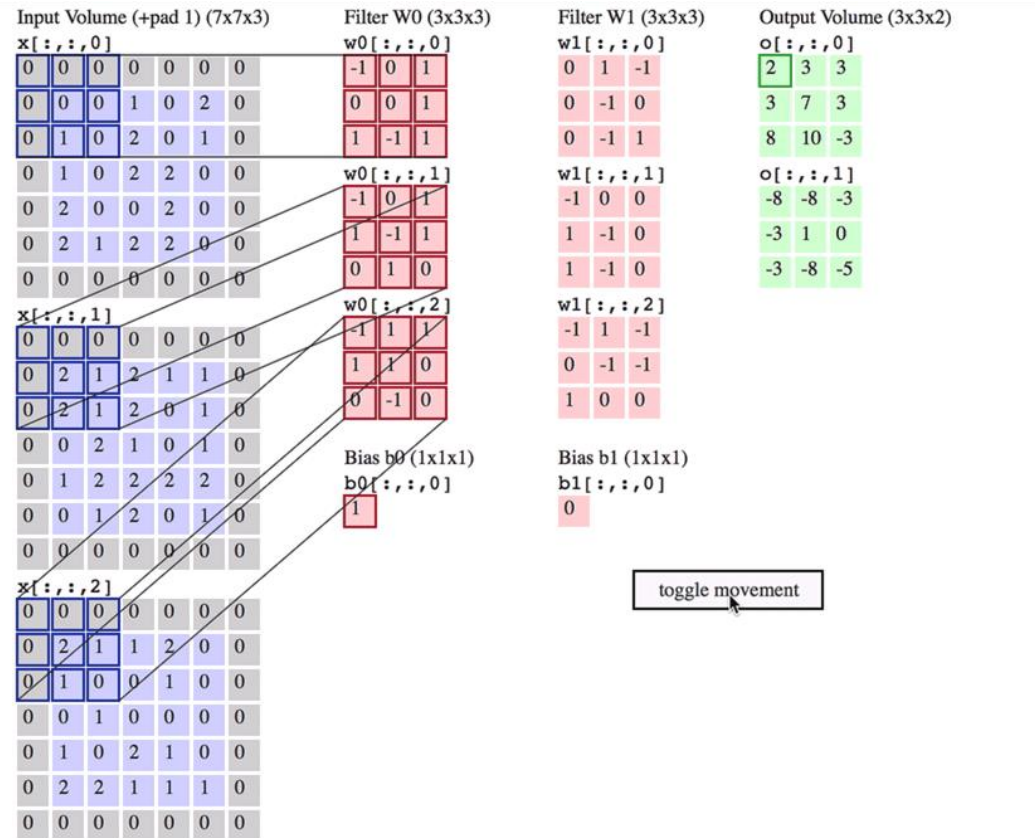
Convolution Operation

- Convolution is using a 'kernel' to extract certain 'features' from an input image.
- A kernel is a matrix, which is slid across the image and multiplied with the input such that the output is enhanced in a certain desirable manner. Kernel is also called a filter.
- Images can be Black/white, Grayscale and RGB (3 channels)
- The output of convolutional layers are feature maps

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:



Convolutional layer – multiple filters example



Feature Extraction with Convolution - Creating Feature Maps



Original



Sharpen




Edge Detect

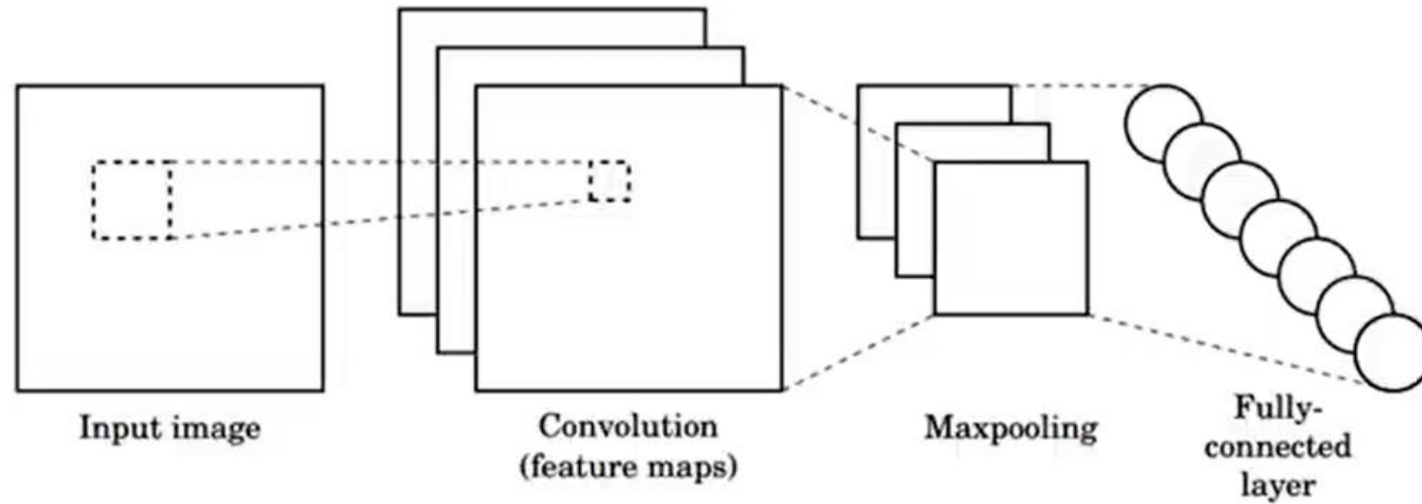


"Strong" Edge
Detect

- Apply a set of weights (Filter) to extract local features
- Different types of Filters are used to extract features.
- Spatially share parameters (weights) of each filter in a convolution layer
- Reducing number of connections between layers




CNNs for Classification



1. **Convolution:** Apply filters to generate feature maps.
2. **Non-linearity:** Often ReLU.
3. **Pooling:** Downsampling operation on each feature map.


 `tf.keras.layers.Conv2D`


 `tf.keras.activations.*`

 `tf.keras.layers.MaxPool2D`

Train model with image data.
Learn weights of filters in convolutional layers.

Convolution Layers: Local Connectivity



 `tf.keras.layers.Conv2D`

For a neuron in hidden layer:

- Take inputs from patch
- Compute weighted sum
- Apply bias

4x4 filter: matrix
of weights w_{ij}

$$\sum_{i=1}^4 \sum_{j=1}^4 w_{ij} x_{i+p,j+q} + b$$


for neuron (p,q) in hidden layer

- 1) applying a window of weights
- 2) computing linear combinations
- 3) activating with non-linear function

Note: Rectified Linear Unit (ReLU) is used after every Convolution operation to replace all negative pixel values in the feature map by zero. tanh or sigmoid can also be used instead of ReLU, but ReLU has been found to perform better in most situations.

MNIST Handwritten Digit Recognition

0–9 handwritten digit recognition:



MNIST Data maintained by Yann LeCun: <http://yann.lecun.com/exdb/mnist/>
Keras provides data sets loading function at <http://keras.io/datasets>


State of the Art Object Classification (Handwritten digit recognition) Results




mnist-cnn-example.py

Pooling Layer

- Pooling layer down samples the volume spatially, independently in each depth slice of the input.
- Popular pooling is Max pooling. Mean pooling can also be used.

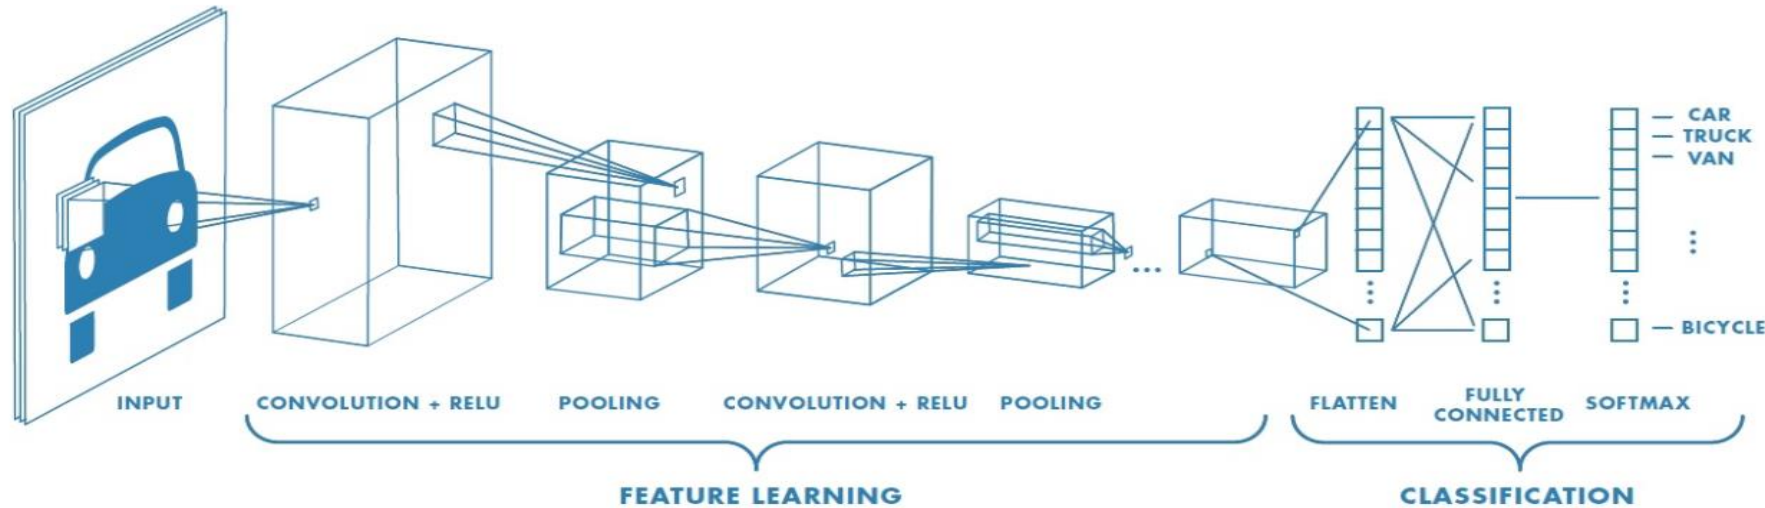


Deep CNNs – Representation Learning



Note: The further you advance into the Deep CNN, the more complex the features your nodes can recognize, since they aggregate and recombine features from the previous layer

CNNs for Classification: Feature Learning + Classification probabilities



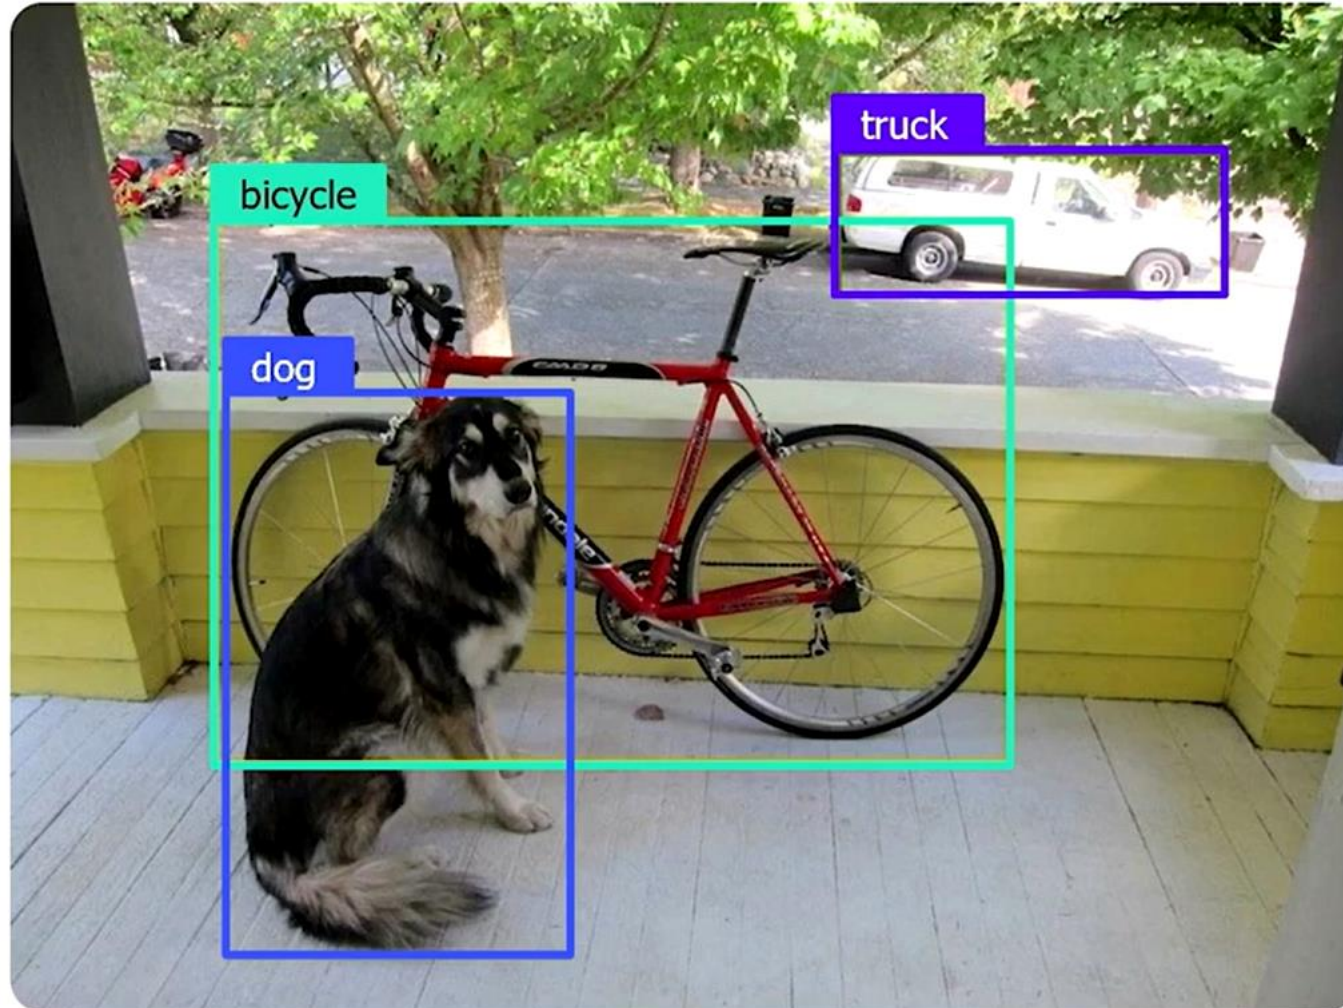
Feature Learning

- Feature Learning in image with convolution
- Non-linearity using activation functions
- Reduce dimensionality and preserve spatial invariance with Pooling
- Convolution and pooling layers extract high level input features.

Classification

- ❖ Fully connected layer uses the extracted features from Feature Learning pipeline.
- ❖ Express output as probability of image belonging to a particular class with the help of SoftMax function.

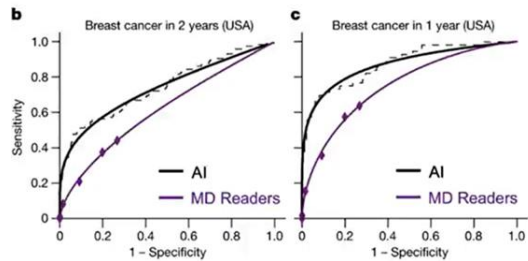
Object Detection using a Deep learning model (e.g., YOLO)




CNN Architecture: Many applications

Classification: Breast Cancer Screening

International evaluation of an AI system for breast cancer screening



CNN-based system outperformed expert radiologists at detecting breast cancer from mammograms




Breast cancer case missed by radiologist but detected by AI

Object Detection



Semantic Segmentation: Fully Convolutional Networks

FCN: Fully Convolutional Network.
Network designed with all convolutional layers,
with **downsampling** and **upsampling** operations



Continuous Control: Navigation from Vision

