


# CSE2008: Operating Systems

L15 L16 L17 & L18: Process Synchronization



Dr. Subrata Tikadar  
SCOPE, VIT-AP University

# Recap

- Introductory Concepts
- Process Fundamentals
- IPC
- CPU Scheduling Algorithms
- Multithreading Concepts

# Outline

- Process synchronization
- Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Mutex Locks
- Semaphores
- Monitors
- Classic problems of Synchronization

# Process synchronization

- Preliminary Concept
  - Producer–Consumer Problem

```
while (true) {
    /* produce an item in next_produced */

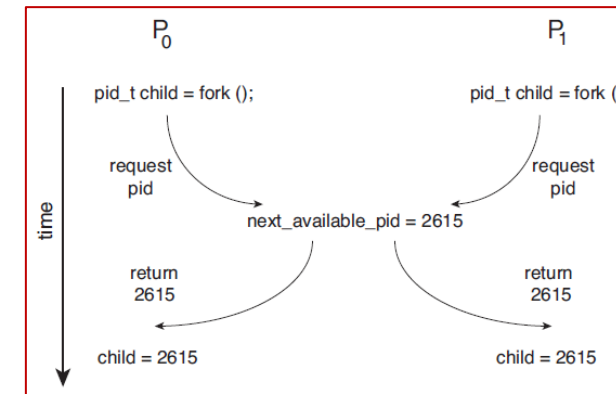
    while (count == BUFFER_SIZE)
        ; /* do nothing */

    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    count++;
}
```

```
register1 = count
register1 = register1 + 1
count = register1
```

T <sub>0</sub> :	producer	execute	register <sub>1</sub> = count	{register <sub>1</sub> = 5}
T <sub>1</sub> :	producer	execute	register <sub>1</sub> = register <sub>1</sub> + 1	{register <sub>1</sub> = 6}
T <sub>2</sub> :	consumer	execute	register <sub>2</sub> = count	{register <sub>2</sub> = 5}
T <sub>3</sub> :	consumer	execute	register <sub>2</sub> = register <sub>2</sub> - 1	{register <sub>2</sub> = 4}
T <sub>4</sub> :	producer	execute	count = register <sub>1</sub>	{count = 6}
T <sub>5</sub> :	consumer	execute	count = register <sub>2</sub>	{count = 4}

```
register2 = count
register2 = register2 - 1
count = register2
```




```
while (true) {
    while (count == 0)
        ; /* do nothing */

    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;

    /* consume the item in next_consumed */
}
```

# Critical-Section Problem

- Critical-Section Problem
  - Critical-Section: A segment of code in which the process may be accessing — and updating — data that is shared with at least one other process



# Critical-Section Problem

- Critical-Section Problem

- A solution to the critical-section problem must satisfy the following three requirements:

1. **Mutual exclusion** → If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections.
2. **Progress** → If no process is executing in its critical section and some processes wish to enter their critical sections, then only those processes that are not executing in their remainder sections can participate in deciding which will enter its critical section next, and this selection cannot be postponed indefinitely
3. **Bounded waiting** → There exists a bound, or limit, on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.

# Peterson's Solution

- Peterson's Solution
  - Restricted to two processes that alternate execution between their critical sections and remainder sections
    - The processes are numbered  $P_0$  and  $P_1$ . For convenience, when presenting  $P_i$ , we use  $P_j$  to denote the other process; that is,  $j$  equals  $1 - i$
  - Requires the two processes to share two data items:

```
int turn;
```

```
boolean flag[2];
```

---

```
while (true) {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j)  
        ;  
  
    /* critical section */  
  
    flag[i] = false;  
  
    /*remainder section */  
}
```

# Peterson's Solution

- Peterson's Solution
  - To prove that this solution is correct, we need to show that:
    1. Mutual exclusion is preserved,
    2. The progress requirement is satisfied, and
    3. The bounded-waiting requirement is met.

```
int turn;
boolean flag[2];



---



while (true) {
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j)
        ;

    /* critical section */

    flag[i] = false;

    /*remainder section */
}
```

---



# Peterson's Solution

- Peterson's Solution

- To prove that this solution is correct, we need to show that:

1. Mutual exclusion is preserved,
2. The progress requirement is satisfied, and
3. The bounded-waiting requirement is met.

Each  $P_i$  enters its critical section only if either  $\text{flag}[j] == \text{false}$  or  $\text{turn} == i$ . Also note that, if both processes can be executing in their critical sections at the same time, then  $\text{flag}[0] == \text{flag}[1] == \text{true}$ . These two observations imply that  $P_0$  and  $P_1$  could not have successfully executed their while statements at about the same time, since the value of  $\text{turn}$  can be either 0 or 1 but cannot be both. Hence, one of the processes—say,  $P_j$ —must have successfully executed the while statement, whereas  $P_i$  had to execute at least one additional statement (“ $\text{turn} == j$ ”). However, at that time,  $\text{flag}[j] == \text{true}$  and  $\text{turn} == j$ , and this condition will persist as long as  $P_j$  is in its critical section; as a result, **mutual exclusion** is preserved.

```
int turn;  
boolean flag[2];
```

---

```
while (true) {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j)  
        ;  
  
    /* critical section */  
  
    flag[i] = false;  
  
    /*remainder section */  
}
```

---

# Peterson's Solution

- Peterson's Solution

- To prove that this solution is correct, we need to show that:

1. Mutual exclusion is preserved,
2. The progress requirement is satisfied, and
3. The bounded-waiting requirement is met.

A process  $P_i$  can be prevented from entering the critical section only if it is stuck in the while loop with the condition  $\text{flag}[j] == \text{true}$  and  $\text{turn} == j$ ; this loop is the only one possible. If  $P_j$  is not ready to enter the critical section, then  $\text{flag}[j] == \text{false}$ , and  $P_i$  can enter its critical section. If  $P_j$  has set  $\text{flag}[j]$  to true and is also executing in its while statement, then either  $\text{turn} == i$  or  $\text{turn} == j$ . If  $\text{turn} == i$ , then  $P_i$  will enter the critical section. If  $\text{turn} == j$ , then  $P_j$  will enter the critical section. However, once  $P_j$  exits its critical section, it will reset  $\text{flag}[j]$  to false, allowing  $P_i$  to enter its critical section. If  $P_j$  resets  $\text{flag}[j]$  to true, it must also set  $\text{turn}$  to  $i$ . Thus, since  $P_i$  does not change the value of the variable  $\text{turn}$  while executing the while statement,  $P_i$  will enter the critical section (**progress**) after at most one entry by  $P_j$  (**bounded waiting**).

```
int turn;  
boolean flag[2];
```

---

```
while (true) {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j)  
        ;  
  
    /* critical section */  
  
    flag[i] = false;  
  
    /*remainder section */  
}
```

---

# Synchronization Hardware

- Memory Barriers
- Hardware Instructions
- Atomic Variables

# Synchronization Hardware

- Memory Barriers

- Memory Model:

1. Strongly ordered → where a memory modification on one processor is immediately visible to all other processors
2. Weakly ordered → where modifications to memory on one processor may not be immediately visible to other processors

➤ Instructions that can force any changes in memory to be propagated to all other processors, thereby ensuring that memory modifications are visible to threads running on other processors – such instructions are known as **memory barriers** or **memory fences**.

# Synchronization Hardware

- Hardware Instructions
  - test and set() instruction – TAS

---

```
boolean test_and_set(boolean *target) {  
    boolean rv = *target;  
    *target = true;  
  
    return rv;  
}
```

---

---

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = false;  
  
    /* remainder section */  
} while (true);
```

---

# Synchronization Hardware

- Hardware Instructions
  - compare\_and\_swap() instruction – CAS

---

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
  
    return temp;  
}
```

---

---

```
while (true) {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
}
```

---

# Synchronization Hardware

- Hardware Instructions
  - compare\_and\_swap() instruction – CAS

***Self Study:*** Analyse how this algorithm satisfies all the critical-section requirements!

```
while (true) {
    waiting[i] = true;
    key = 1;
    while (waiting[i] && key == 1)
        key = compare_and_swap(&lock,0,1);
    waiting[i] = false;

    /* critical section */

    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;

    if (j == i)
        lock = 0;
    else
        waiting[j] = false;

    /* remainder section */
}
```

# Synchronization Hardware

- Atomic Variables

```
increment(&sequence);
```

where the `increment()` function is implemented using the CAS instruction:

```
void increment(atomic_int *v)
{
    int temp;

    do {
        temp = *v;
    }
    while (temp != compare_and_swap(v, temp, temp+1));
}
```



# Mutex Locks

---

```
while (true) {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
}
```

---

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}
```

```
release() {  
    available = true;  
}
```

# Semaphores

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

# Semaphores

- Categories of Semaphores

- Counting semaphore

- The value of it can range over an unrestricted domain

- Binary semaphore

- The value of a binary semaphore can range only between 0 and 1

# Semaphores

- Semaphore Usage

- Example:

- Let us assume that two concurrently running processes:  $P_1$  with a statement  $S_1$  and  $P_2$  with a statement  $S_2$
    - Also assume that we require that  $S_2$  be executed only after  $S_1$  has completed

```
 $S_1$  ;  
signal(synch) ;
```

```
wait(synch) ;  
 $S_2$  ;
```

# Semaphores

- Semaphore Implementation

- Modified definitions of the wait() and signal() operations [to avoid busy waiting]

*When a process executes the wait() operation and finds that the semaphore value is not positive, it must wait. However, rather than engaging in busy waiting, the process can suspend itself. The suspend operation places a process into a waiting queue associated with the semaphore, and the state of the process is switched to the waiting state. Then control is transferred to the CPU scheduler, which selects another process to execute.*

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        sleep();  
    }  
}
```

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

# Semaphores

- Pros and Cons

- Pros – provides a convenient and effective mechanism for process synchronization
- Cons – incorrect usage can result in timing errors that are difficult to detect


- Example:

Suppose that a program interchanges the order in which the wait() and signal() operations on the semaphore mutex are executed, resulting in the following execution:

<code>signal(mutex);</code>	<code>wait(mutex);</code>
<code>...</code>	<code>...</code>
<code>critical section</code>	<code>critical section</code>
<code>...</code>	<code>...</code>
<code>wait(mutex);</code>	<code>wait(mutex);</code>

# Monitors

- An ADT (abstract data type) that includes a set of programmer-defined operations that are provided with mutual exclusion within the monitor
- Also declares the variables whose values define the state of an instance of that type, along with the bodies of functions that operate on those variables



```
monitor monitor name
{
    /* shared variable declarations */

    function P1 ( . . . ) {
        . . .
    }

    function P2 ( . . . ) {
        . . .
    }

    .
    .
    .
    function Pn ( . . . ) {
        . . .
    }

    initialization_code ( . . . ) {
        . . .
    }
}
```


# Monitors

- The condition construct
  - A programmer who needs to write a tailor-made synchronization scheme can define one or more variables of type condition:

condition x, y;

*The only operations that can be invoked on a condition variable are **wait()** and **signal()**.*

*The operation **x.wait()**; means that the process invoking this operation is suspended until another process invokes **x.signal()**;*





# Monitors

- The condition construct
  - A programmer who needs to write a tailor-made synchronization scheme can define one or more variables of type condition:

condition x, y;

*The only operations that can be invoked on a condition variable are **wait()** and **signal()**.*

*The operation **x.wait()**; means that the process invoking this operation is suspended until another process invokes **x.signal()**;*

Now suppose that, when the x.signal() operation is invoked by a process P, there exists a suspended process Q associated with condition x.

→ Two possibilities exist:

1. **Signal and wait.** P either waits until Q leaves the monitor or waits for another condition.
2. **Signal and continue.** Q either waits until P leaves the monitor or waits for another condition.

# Classic problems of Synchronization

- The Bounded-Buffer Problem
- The Readers–Writers Problem
- The Dining-Philosophers Problem

# Classic problems of Synchronization

- The Bounded-Buffer Problem

- Let us assume that the producer and consumer processes share the following data structures:

```
int n;  
semaphore mutex = 1;  
semaphore empty = n;  
semaphore full = 0
```

---

```
while (true) {  
    . . .  
    /* produce an item in next_produced */  
    . . .  
    wait(empty);  
    wait(mutex);  
    . . .  
    /* add next_produced to the buffer */  
    . . .  
    signal(mutex);  
    signal(full);  
}
```

---

```
while (true) {  
    wait(full);  
    wait(mutex);  
    . . .  
    /* remove an item from buffer to next_consumed */  
    . . .  
    signal(mutex);  
    signal(empty);  
    . . .  
    /* consume the item in next_consumed */  
    . . .  
}
```

---

# Classic problems of Synchronization

- The Readers–Writers Problem

- Readers – the processes that want only to read the database
- Writers – the processes that want to update the database

If at least one process among the set of concurrent processes is a *writer* then chaos may occur

➤ To ensure that these difficulties do not arise, we require that the writers have exclusive access to the shared database while writing to the database

*This synchronization problem is referred to as **the readers–writers problem**.*

- **The first readers–writers problem** – *No reader be kept waiting unless a writer has already obtained permission to use the shared object*
- **The second readers–writers problem** – *If a writer is waiting to access the object, no new readers may start reading*

# Classic problems of Synchronization

- The Readers–Writers Problem

- Solution to the first readers-writers problem

- The reader processes share the following data structures:

```
semaphore rw_mutex = 1;  
semaphore mutex = 1;  
int read_count = 0;
```

- Codes for a writer process and a reader process:

---

```
while (true) {  
    wait(rw_mutex);  
    . . .  
    /* writing is performed */  
    . . .  
    signal(rw_mutex);  
}
```

---


---

```
while (true) {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
    . . .  
    /* reading is performed */  
    . . .  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
}
```

---

# Classic problems of Synchronization

- The Dining-Philosophers Problem




## Semaphore Solution

- The shared data:  
semaphore chopstick[5];
- The structure of philosopher i:

```
while (true) {  
    wait(chopstick[i]);  
    wait(chopstick[(i+1) % 5]);  
    . . .  
    /* eat for a while */  
    . . .  
    signal(chopstick[i]);  
    signal(chopstick[(i+1) % 5]);  
    . . .  
    /* think for awhile */  
    . . .  
}
```

# Classic problems of Synchronization

- The Dining-Philosophers Problem



## Monitor Solution

- The data structure:  
enum {THINKING, HUNGRY, EATING} state[5];  
condition self[5];
- The structure of philosopher i:

```
monitor DiningPhilosophers
{
    enum {THINKING, HUNGRY, EATING} state[5];
    condition self[5];

    void pickup(int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self[i].wait();
    }


    void putdown(int i) {
        state[i] = THINKING;
        test((i + 4) % 5);
        test((i + 1) % 5);
    }

    void test(int i) {
        if ((state[(i + 4) % 5] != EATING) &&
            (state[i] == HUNGRY) &&
            (state[(i + 1) % 5] != EATING)) {
            state[i] = EATING;
            self[i].signal();
        }
    }

    initialization_code() {
        for (int i = 0; i < 5; i++)
            state[i] = THINKING;
    }
}
```

# Classic problems of :

- The Dining-Philosophers Problem



Monitor

■

■

```
monitor DiningPhilosophers
{
    enum {THINKING, HUNGRY, EATING} state[5];
    condition self[5];

    void pickup(int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self[i].wait();
    }

    void putdown(int i) {
        state[i] = THINKING;
        test((i + 4) % 5);
        test((i + 1) % 5);
    }


    void test(int i) {
        if ((state[(i + 4) % 5] != EATING) &&
            (state[i] == HUNGRY) &&
            (state[(i + 1) % 5] != EATING)) {
            state[i] = EATING;
            self[i].signal();
        }
    }

    initialization_code() {
        for (int i = 0; i < 5; i++)
            state[i] = THINKING;
    }
}
```



# Classic problems of Synchronization

- The Dining-Philosophers Problem



## Monitor Solution

- The data structure:  
enum {THINKING, HUNGRY, EATING} state[5];  
condition self[5];
- The Philosopher i must invoke the operations pickup() and putdown() in the following sequence:

```
DiningPhilosophers.pickup(i);
```

```
...
```

```
eat
```

```
...
```

```
DiningPhilosophers.putdown(i);
```

# Reference

- Abraham Silberschatz , Peter B. Galvin, Greg Gagne, “Operating System Concepts”, Addison Wesley, 10th edition, 2018
  - Chapter 6: Section 6.1 – 6.7
  - Chapter 7: Section 7.1 – 7.2

# Next

- Deadlock

# Thank You